

```
In [ ]: pip install --upgrade xgboost
```

```
Requirement already satisfied: xgboost in /usr/local/lib/python3.10/dist-pack
ages (2.1.2)
Collecting xgboost
  Downloading xgboost-2.1.3-py3-none-manylinux_2_28_x86_64.whl.metadata (2.1
kB)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packag
es (from xgboost) (1.26.4)
Requirement already satisfied: nvidia-nccl-cu12 in /usr/local/lib/python3.10/
dist-packages (from xgboost) (2.23.4)
Requirement already satisfied: scipy in /usr/local/lib/python3.10/dist-packag
es (from xgboost) (1.13.1)
Downloading xgboost-2.1.3-py3-none-manylinux_2_28_x86_64.whl (153.9 MB)
_____ 153.9/153.9 MB 5.5 MB/s eta 0:00:
00
Installing collected packages: xgboost
  Attempting uninstall: xgboost
    Found existing installation: xgboost 2.1.2
    Uninstalling xgboost-2.1.2:
      Successfully uninstalled xgboost-2.1.2
  Successfully installed xgboost-2.1.3
```

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, OneHotEncoder, StandardScaler
from sklearn.linear_model import LinearRegression, Lasso, Ridge
from sklearn.model_selection import train_test_split, cross_val_score, GridSea
rchCV, KFold, RandomizedSearchCV
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, ExtraTreesRegressor, Gradi
entBoostingRegressor
import xgboost as xgb
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
from yellowbrick.regressor import PredictionError, ResidualsPlot
import os
import warnings
warnings.filterwarnings("ignore")

pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', None)
```

Read dataset

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')

data = pd.read_csv('/content/drive/My Drive/ML/dataset/AB_NYC_2019.csv')
data.head()
```

Mounted at /content/drive

Out[]:

	id	name	host_id	host_name	neighbourhood_group	neighbourhood	latitude
0	2539	Clean & quiet apt home by the park	2787	John	Brooklyn	Kensington	40.64749
1	2595	Skylit Midtown Castle	2845	Jennifer	Manhattan	Midtown	40.75362
2	3647	THE VILLAGE OF HARLEM....NEW YORK !	4632	Elisabeth	Manhattan	Harlem	40.80902
3	3831	Cozy Entire Floor of Brownstone	4869	LisaRoxanne	Brooklyn	Clinton Hill	40.68514
4	5022	Entire Apt: Spacious Studio/Loft by central park	7192	Laura	Manhattan	East Harlem	40.79851



EDA and Data Visualisations

Exploratory Data Analysis was performed on the dataset to have a better understanding of the dataset.

In []: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 48895 entries, 0 to 48894
Data columns (total 16 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    48895 non-null  int64
1   name                                48879 non-null  object
2   host_id                             48895 non-null  int64
3   host_name                           48874 non-null  object
4   neighbourhood_group                 48895 non-null  object
5   neighbourhood                       48895 non-null  object
6   latitude                           48895 non-null  float64
7   longitude                           48895 non-null  float64
8   room_type                           48895 non-null  object
9   price                               48895 non-null  int64
10  minimum_nights                      48895 non-null  int64
11  number_of_reviews                   48895 non-null  int64
12  last_review                         38843 non-null  object
13  reviews_per_month                  38843 non-null  float64
14  calculated_host_listings_count      48895 non-null  int64
15  availability_365                    48895 non-null  int64
dtypes: float64(3), int64(7), object(6)
memory usage: 6.0+ MB
```

Room type analysis

In []: `room_type = data.groupby('room_type').agg(avg_price=('price', 'mean'),
count=('price', 'count')).reset_index()`

room_type

Out[]:

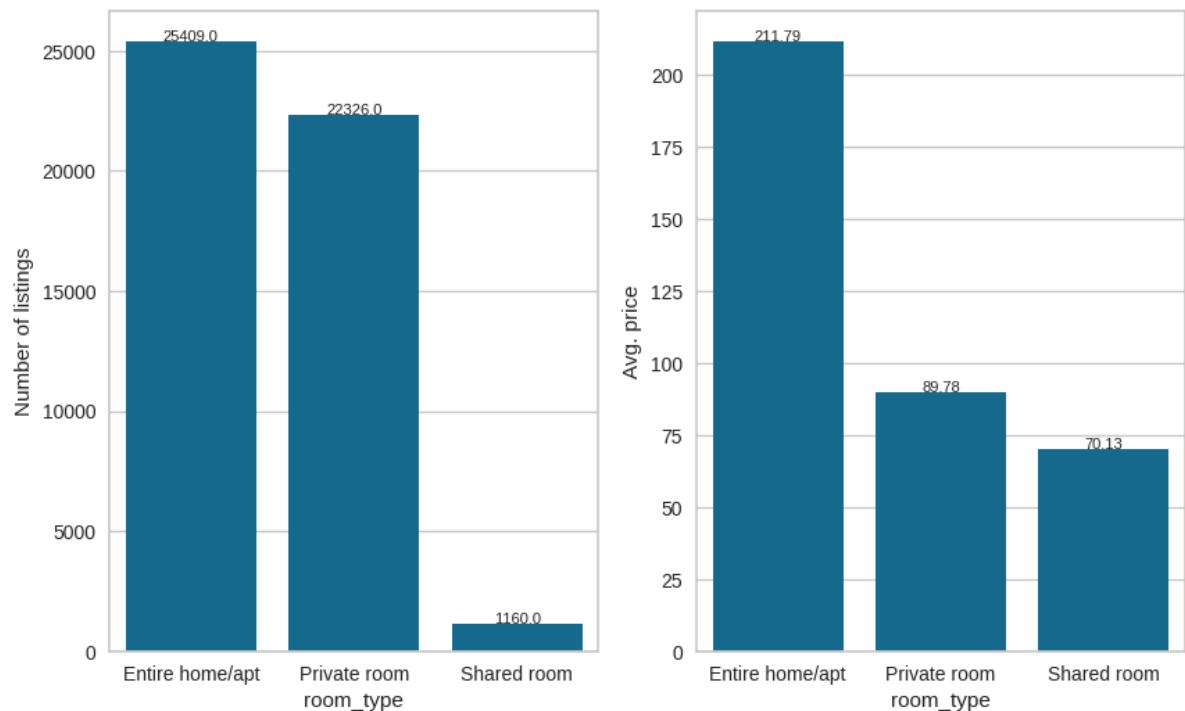
	room_type	avg_price	count
0	Entire home/apt	211.794246	25409
1	Private room	89.780973	22326
2	Shared room	70.127586	1160

```

In [ ]: fig,(chart_1,chart_2) = plt.subplots(1,2,figsize=(10,6))
ax = sns.barplot(x='room_type', y = 'avg_price', data=room_type, ax=chart_2)
for p in chart_2.patches:
    chart_2.annotate("{:.2f}".format(p.get_height()),
                     (p.get_x() + p.get_width()/2., p.get_height()),
                     ha='center',
                     fontsize=8)
chart_2.set_ylabel("Avg. price")

ax = sns.barplot(x='room_type', y = 'count', data=room_type, ax=chart_1)
for p in chart_1.patches:
    chart_1.annotate("{:.1f}".format(p.get_height()),
                     (p.get_x() + p.get_width()/2., p.get_height()),
                     ha='center',
                     fontsize=8)
chart_1.set_ylabel("Number of listings")
plt.show()

```



The dataset has 3 different room types:

- 52% (entire home/apartment)
- 45% (private room)
- 3% (shared room)

Also, Entire home/apt are most expensive.

Neighbourhood_group(large area) analysis

```
In [ ]: # avg price of neighbourhoods
neighbourhood_group_df = data.groupby('neighbourhood_group').agg(avg_price=('p
rice', 'mean'),
                                                                    count=('price', 'count')).reset_index()

neighbourhood_group_df
```


Out[]:

	neighbourhood_group	avg_price	count
0	Bronx	87.496792	1091
1	Brooklyn	124.383207	20104
2	Manhattan	196.875814	21661
3	Queens	99.517649	5666
4	Staten Island	114.812332	373

```
In [ ]: #Pie chart to show Listings in each neighbourhood group
ax1 = data.neighbourhood_group.value_counts().plot.pie(title='Neighborhood Groups',
figsize=(6,6),
autopct='%1.1f%%')

ax1.set_ylabel(None)
```

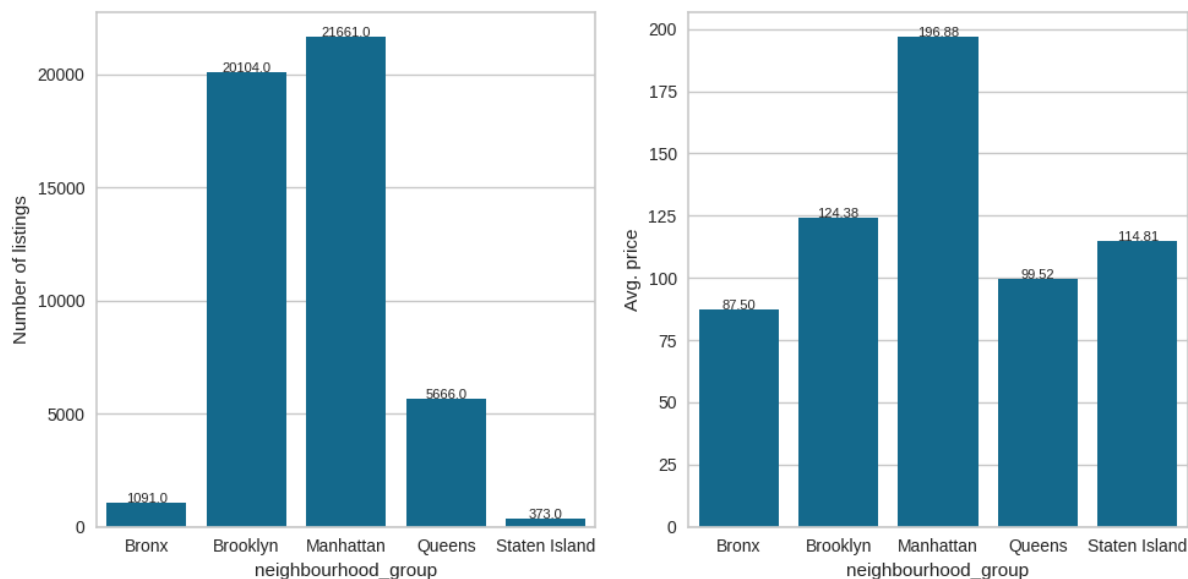
```
Out[ ]: Text(0, 0.5, '')
```



More than 85% of listings are located in Manhattan and Brooklyn.

```
In [ ]: fig,(chart_1,chart_2) = plt.subplots(1,2,figsize=(10,5))
ax = sns.barplot(x='neighbourhood_group', y = 'avg_price', data=neighbourhood_group_df, ax=chart_2)
for p in chart_2.patches:
    chart_2.annotate("{:.2f}".format(p.get_height()),
                      (p.get_x() + p.get_width()/2., p.get_height()),
                      ha='center',
                      fontsize=8)
chart_2.set_ylabel("Avg. price")

ax = sns.barplot(x='neighbourhood_group', y = 'count', data=neighbourhood_group_df, ax=chart_1)
for p in chart_1.patches:
    chart_1.annotate("{:.1f}".format(p.get_height()),
                     (p.get_x() + p.get_width()/2., p.get_height()),
                     ha='center',
                     fontsize=8)
chart_1.set_ylabel("Number of listings")
plt.tight_layout()
plt.show()
```



The pricing for the properties was skewed highly towards Manhattan at \$196, while all the other locations fell in the \$85-\$125 price range.

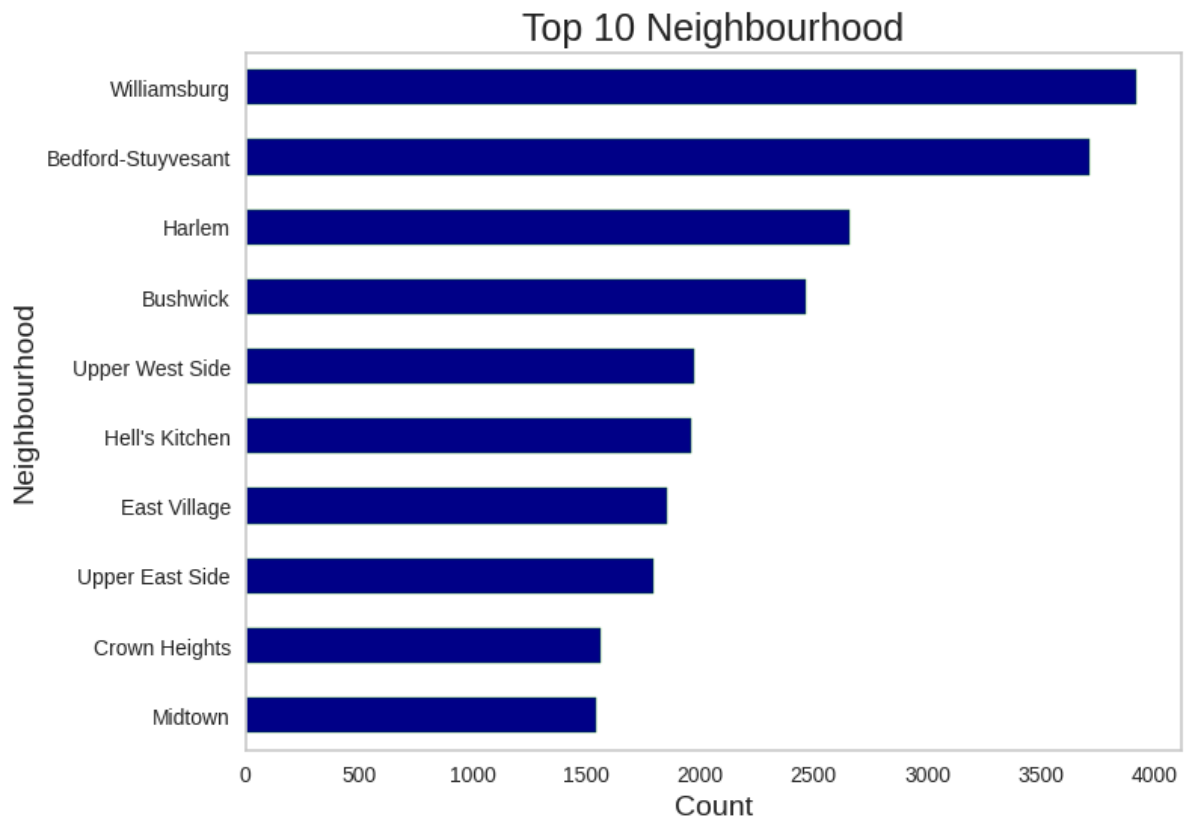
Neighbourhood Analysis

```
In [ ]: print("There are", data["neighbourhood"].nunique(), "distinct values.")
```

There are 221 distinct values.

```
In [ ]: plt.figure(figsize=(8,6))
data["neighbourhood"].value_counts().head(10).sort_values().plot(kind = "bar",
                                                                    color = "darkblue",
                                                                    edgecolor="green",
                                                                    grid =False)
ax.bar_label(ax.containers[0], fmt='%.0f', color="black", fontsize=10)
plt.xlabel("Count", size = 14)
plt.ylabel("Neighbourhood", size = 14)
plt.title("Top 10 Neighbourhood", size = 18)
```

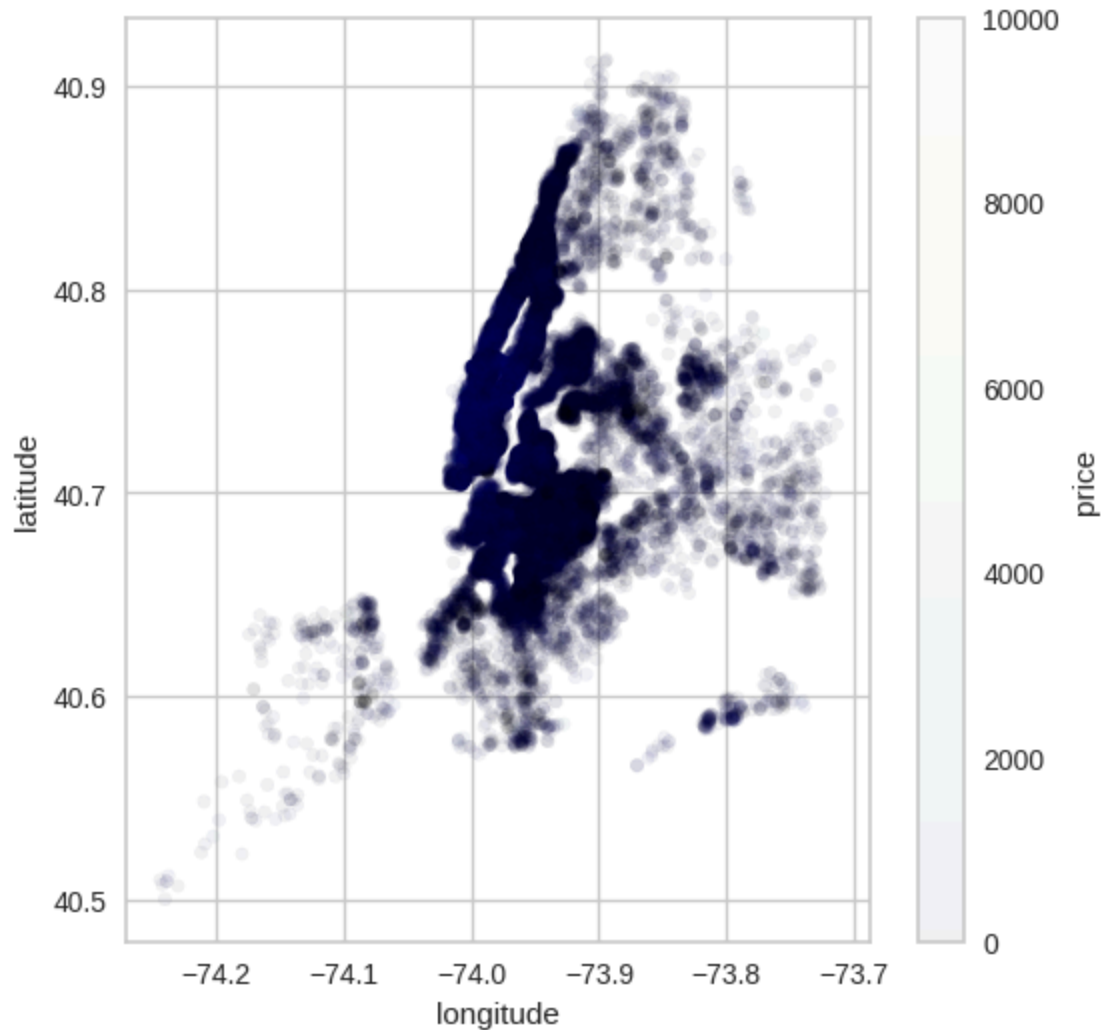
```
Out[ ]: Text(0.5, 1.0, 'Top 10 Neighbourhood')
```



neighbourhood_group and room type

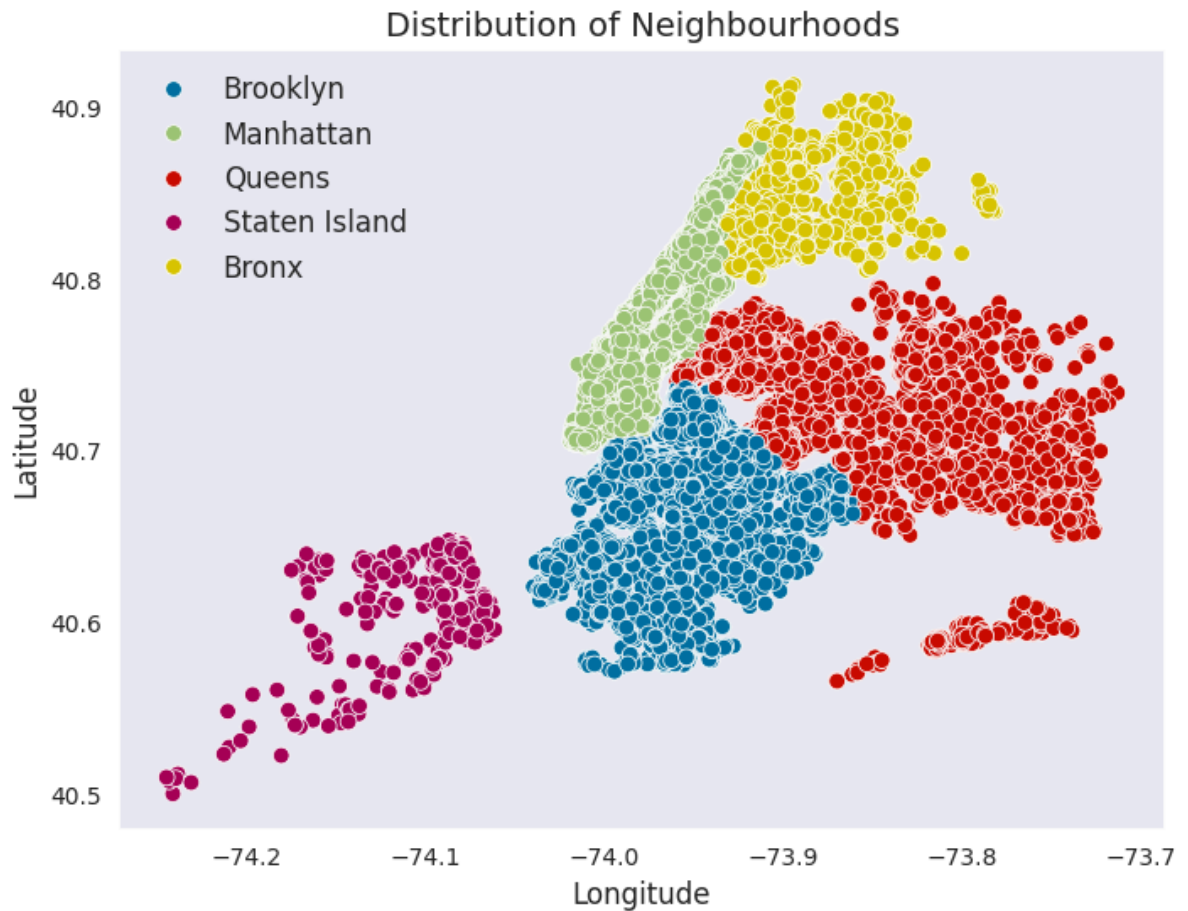
```
In [ ]: data.plot(kind = "scatter", x = "longitude", y = "latitude", alpha = 0.05,  
                figsize = (6,6), c = "price", cmap = plt.get_cmap("gist_earth"),  
                colorbar = True  
                )
```

```
Out[ ]: <Axes: xlabel='longitude', ylabel='latitude'>
```




```
In [ ]: plt.figure(figsize=(8,6))
sns.set_style("dark", {'axes.grid' : False})
sns.scatterplot(data=data, x="longitude", y="latitude", hue="neighbourhood_group")
plt.ylabel("Latitude", fontsize=12)
plt.xlabel("Longitude", fontsize=12)
plt.title("Distribution of Neighbourhoods", fontsize=14)
plt.legend(prop={"size":12})
```

Out[]: <matplotlib.legend.Legend at 0x7b97b4f043d0>



```
In [ ]: plt.figure(figsize=(8,6))
sns.scatterplot(data = data, x="longitude", y="latitude", hue="room_type")
plt.ylabel("Latitude", fontsize=12)
plt.xlabel("Longitude", fontsize=12)
plt.title("Distribution of Room Type", fontsize=14)
plt.legend(prop={"size":12})
```

Out[]: <matplotlib.legend.Legend at 0x7b97a0b9a0b0>



Price Distribution analysis

```
In [ ]: # Log transformation
fig, ax = plt.subplots(1,2, figsize = (10,5))
sns.histplot(data=data, x="price", kde=True, ax=ax[0]).set_title("Price Distri
bution Before Log Transformation",
                                                                size = 12)
sns.histplot(np.log1p(data.price), kde=True, ax=ax[1]).set_title("Price Distri
bution After Log Transformation",
                                                                size = 12)
```

```
Out[ ]: Text(0.5, 1.0, 'Price Distribution After Log Transformation')
```

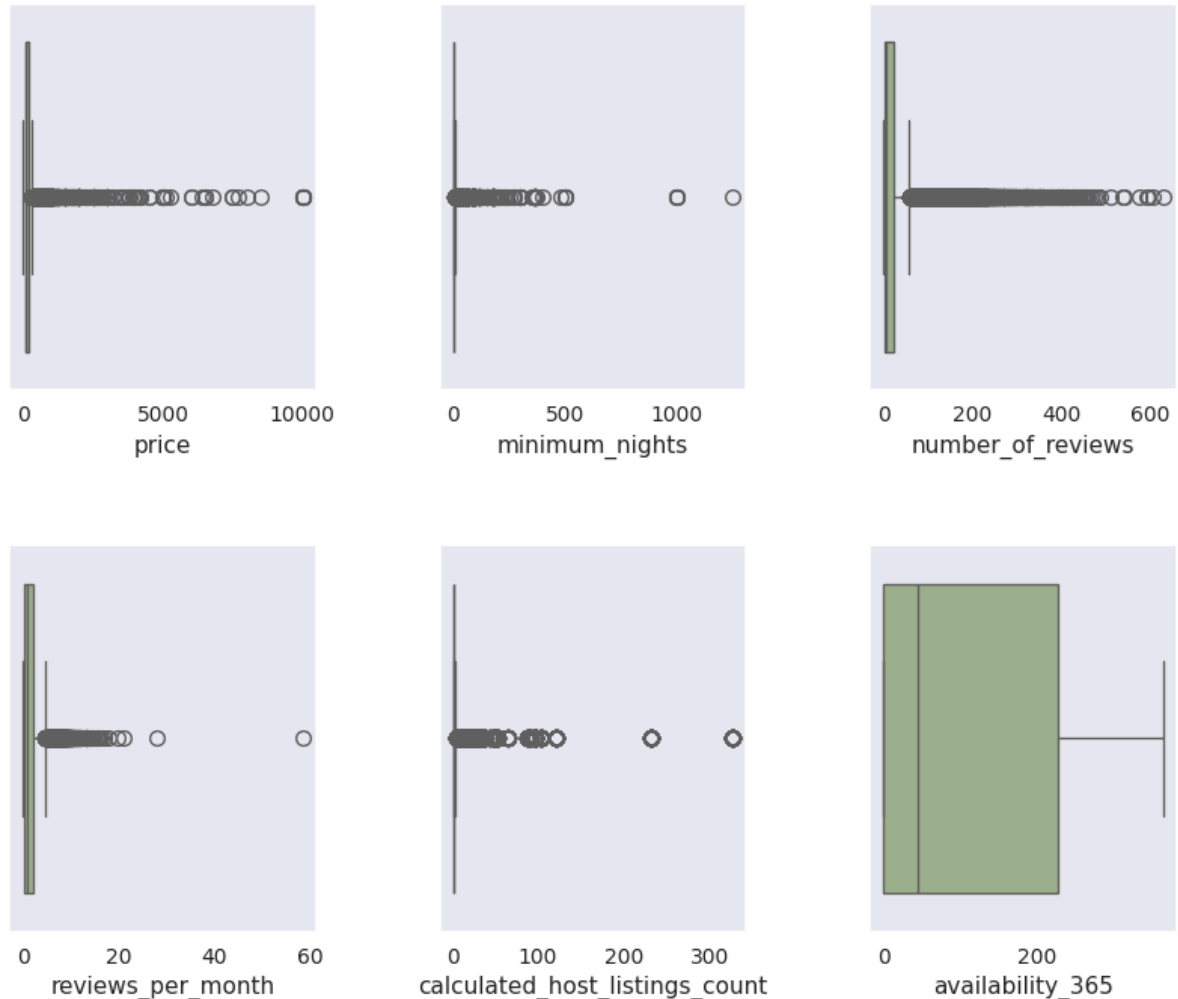


The graph on the left shows that there is a right skewed distribution. To make a better statistical analysis and to get better scores, we will be applied log transformation for the price column

Data Cleaning and Preparation

Outlier detection and Removal

```
In [ ]: columns = ["price", "minimum_nights", "number_of_reviews", "reviews_per_mont
h", "calculated_host_listings_count",
                  "availability_365"]
fig = plt.figure(figsize=(10,8))
fig.subplots_adjust(hspace=0.4, wspace=0.4)
for num, column_name in enumerate(columns):
    ax = fig.add_subplot(2, 3, num + 1)
    ax = sns.boxplot(x=data[column_name], color='#9db787')
```



```
In [ ]: q1_price = data["price"].quantile(0.25)
q3_price = data["price"].quantile(0.75)
iqr_price = q3_price - q1_price
lower_limit_price = q1_price - 1.5 * iqr_price
upper_limit_price = q3_price + 1.5 * iqr_price

df_filter_price = data[(data["price"] > lower_limit_price) & (data["price"] <
upper_limit_price)]
```

```
In [ ]: fig = plt.figure(figsize=(7,4))
fig.subplots_adjust(hspace=0.4, wspace=0.4)

ax = fig.add_subplot(1, 1, 1)
sns.boxplot(x=df_filter_price["price"], color='#9db787')
ax.set_title("Price Distribution")
```

```
Out[ ]: Text(0.5, 1.0, 'Price Distribution')
```



```
In [ ]: df = df_filter_price
```

Correlation matrix


```
In [ ]: # Check data types of all columns
print("Numeric columns:", df.select_dtypes(include=[np.number]).columns.tolist())
print("\nNon-numeric columns:", df.select_dtypes(exclude=[np.number]).columns.tolist())
```

Numeric columns: ['id', 'host_id', 'latitude', 'longitude', 'price', 'minimum_nights', 'number_of_reviews', 'reviews_per_month', 'calculated_host_listings_count', 'availability_365']

Non-numeric columns: ['name', 'host_name', 'neighbourhood_group', 'neighbourhood', 'room_type', 'last_review']

```
In [ ]: # Get numeric columns only
numeric_df = df.select_dtypes(include=[np.number])

# Create the correlation heatmap
plt.figure(figsize=(10,8))
df_corr = numeric_df.corr()
mask = np.triu(np.ones_like(df_corr, dtype=bool))
sns.heatmap(df_corr, mask=mask, annot=True, cmap="GnBu")
plt.tight_layout()
plt.show()
```



Handling Missing Values

```
In [ ]: data.isnull().sum()
```

Out[]:

	0
id	0
name	16
host_id	0
host_name	21
neighbourhood_group	0
neighbourhood	0
latitude	0
longitude	0
room_type	0
price	0
minimum_nights	0
number_of_reviews	0
last_review	10052
reviews_per_month	10052
calculated_host_listings_count	0
availability_365	0

dtype: int64

```
In [ ]: data[data['last_review'].isnull()][['number_of_reviews', 'reviews_per_month']].head()
```

Out[]:

	number_of_reviews	reviews_per_month
2	0	NaN
19	0	NaN
26	0	NaN
36	0	NaN
38	0	NaN


```
In [ ]: # filling non-values of 'reviews_per_month' with 0
data.fillna({'reviews_per_month':0}, inplace=True)
# Drop 'last_review' feature
data.drop('last_review', inplace=True, axis=1)
data.isnull().sum()
```

Out[]:

	0
id	0
name	16
host_id	0
host_name	21
neighbourhood_group	0
neighbourhood	0
latitude	0
longitude	0
room_type	0
price	0
minimum_nights	0
number_of_reviews	0
reviews_per_month	0
calculated_host_listings_count	0
availability_365	0

dtype: int64

If a certain listing has a null for its "last_review," that means it has not gotten a review at all(also proved from previous dataframe), so "reviews_per_month" must be 0.

Also, now we can drop the last-review column containing null values from a cleaner dataset.

drop the columns that will not affect the price prediction.

```
In [ ]: data.drop(["id", "name", "host_id", "host_name"], axis = 1, inplace = True)
```

Applying log transformation for the price column.

```
In [ ]: data["price"] = np.log1p(data["price"])
X = data.drop("price", axis = 1)
y = data["price"]
```

Splitting data into training and testing sets

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)
X_train = X_train.reset_index(drop=True)
X_test = X_test.reset_index(drop=True)
```

Feature Encoding

Encoding is a vital process where categorical variables are transformed into numerical format. One-hot encoding was adopted on 'room type', 'neighbourhood group', and 'neighbourhood'. It was ensured that multi-collinearity was avoided while encoding. It resulted in a total of 235 features.

```
In [ ]: # Using One Hot Encoding
ohe = OneHotEncoder(handle_unknown = 'ignore')
columns = ["neighbourhood_group", "neighbourhood", "room_type"]
ohe_df_train = pd.DataFrame(ohe.fit_transform(X_train[columns]).toarray(), columns=ohe.get_feature_names_out())
X_train_ohe = X_train.join(ohe_df_train).drop(columns, axis=1)

ohe_df_test = pd.DataFrame(ohe.transform(X_test[columns]).toarray(), columns=ohe.get_feature_names_out())
X_test_ohe = X_test.join(ohe_df_test).drop(columns, axis=1)
```

Feature Scaling

Feature scaling transforms numerical features to have the same range. Here, we utilised the standard scaler from the Scikit-Learn library

```
In [ ]: # scale the features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train_ohe)
X_test = scaler.transform(X_test_ohe)
```

Model Building and Comparison

```
In [ ]: classic_ml_models = []
ensemble_bagging_methods = []
ensemble_boosting_methods = []
```

Function for prediction and evaluation

```
In [ ]: def test_models(X_train, X_test, y_train, y_test, model, model_name, models):
    y_pred = model.predict(X_test)

    test_score = r2_score(y_test, y_pred)
    train_score = model.score(X_train, y_train)
    mae = mean_absolute_error(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)
    rmse = np.sqrt(mean_squared_error(y_test, y_pred))

    rmse_cv = np.sqrt(-cross_val_score(model, X_train, y_train, scoring='neg_mean_squared_error', cv=5).mean())

    row = {"Model": model_name, "Test Sc.": test_score, "Train Sc.": train_score,
          "MAE": mae, "MSE": mse, "RMSE": rmse, "RMSE CV": rmse_cv}

    display(pd.DataFrame([row]))
    models.append(row)
```

Funtion for visualising the results together for comparison

```
In [ ]: # Visualization Fuction
sns.set_style("dark", {'axes.grid' : False})
def compare_models(models, rotation):
    display(models.style.highlight_min(subset=["MAE", "MSE", "RMSE", "RMSE CV"],
color = '#379949', axis = 0)\
                                .highlight_max(subset=["Test Sc."], color = '#379949',
axis = 0)\
                                .highlight_min(subset=["Test Sc.,"], color = '#b4241c', axis = 0)\
                                .highlight_max(subset=["MAE", "MSE", "RMSE", "RMSE CV"],
color = '#b4241c', axis = 0))
```

Classic ML models

Ridge regression

```
In [ ]: ridge = Ridge(alpha=1.0)
ridge.fit(X_train, y_train)
test_models(X_train, X_test, y_train, y_test, ridge, "Ridge Regression", classic_ml_models )
```

	Model	Test Sc.	Train Sc.	MAE	MSE	RMSE	RMSE CV
0	Ridge Regression	0.550181	0.537237	0.334657	0.212804	0.461306	0.478069

Lasso Regression

```
In [ ]: lasso = Lasso(alpha = 0.0001)
lasso.fit(X_train, y_train)
test_models(X_train, X_test, y_train, y_test, lasso, "Lasso Regression", classic_ml_models)
```

	Model	Test Sc.	Train Sc.	MAE	MSE	RMSE	RMSE CV
0	Lasso Regression	0.54993	0.537223	0.334679	0.212923	0.461435	0.477871

Decision Tree

```
In [ ]: # hyperparameter set to intermediate value of 60 to prevent underfitting/ over fitting
dtr = DecisionTreeRegressor(min_samples_leaf=60)
dtr.fit(X_train, y_train)
test_models(X_train, X_test, y_train, y_test, dtr, "Decision Tree", classic_ml_models)
```

	Model	Test Sc.	Train Sc.	MAE	MSE	RMSE	RMSE CV
0	Decision Tree	0.571749	0.600147	0.326018	0.2026	0.450111	0.467157

Comparison of classic ML methods

```
In [ ]: print("Green Highlights the best model, and Red highlights the worst.")
compare_models(pd.DataFrame(classic_ml_models),0)
```

Green Highlights the best model, and Red highlights the worst.

	Model	Test Sc.	Train Sc.	MAE	MSE	RMSE	RMSE CV
0	Ridge Regression	0.550181	0.537237	0.334657	0.212804	0.461306	0.478069
1	Lasso Regression	0.549930	0.537223	0.334679	0.212923	0.461435	0.477871
2	Decision Tree	0.571749	0.600147	0.326018	0.202600	0.450111	0.467157

the Decision tree got the highest accuracy and the least errors

Bagging Ensemble Methods

Random Forest

```
In [ ]: # Baseline random forest regression.
rfr = RandomForestRegressor()
rfr.fit(X_train, y_train)
test_models(X_train, X_test, y_train, y_test, rfr, "Random Forest (Base)", ensemble_bagging_methods)
```

	Model	Test Sc.	Train Sc.	MAE	MSE	RMSE	RMSE CV
0	Random Forest (Base)	0.61025	0.942988	0.306271	0.184386	0.429402	0.447824

Hyper parameter Tuning

```
In [ ]: # Parameters currently in use for baseline model.
from pprint import pprint
print('Parameters currently in use: \n')
pprint(rfr.get_params())
```

Parameters currently in use:

```
{'bootstrap': True,
 'ccp_alpha': 0.0,
 'criterion': 'squared_error',
 'max_depth': None,
 'max_features': 1.0,
 'max_leaf_nodes': None,
 'max_samples': None,
 'min_impurity_decrease': 0.0,
 'min_samples_leaf': 1,
 'min_samples_split': 2,
 'min_weight_fraction_leaf': 0.0,
 'monotonic_cst': None,
 'n_estimators': 100,
 'n_jobs': None,
 'oob_score': False,
 'random_state': None,
 'verbose': 0,
 'warm_start': False}
```

```
In [ ]: param_grid = {
    'n_estimators': [100, 300, 500],
    'max_depth': [10, 30],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2],
    'max_features': ['auto', 'sqrt']
}

rf_random = RandomizedSearchCV(
    estimator=rfr,
    param_distributions=param_grid,
    n_iter=5,
    cv=3,
    verbose=2,
    random_state=42,
    n_jobs=2
)
```

```
In [ ]: rf_random.fit(X_train, y_train)
```

Fitting 3 folds for each of 5 candidates, totalling 15 fits

```
Out[ ]: RandomizedSearchCV
  best_estimator_: RandomForestRegressor
    RandomForestRegressor
```

(https://scikit-learn.org/1.5/modules/generated/sklearn.model_selection.RandomizedSearchCV.html)

(<https://scikit-learn.org/1.5/modules/generated/sklearn.ensemble.RandomForestRegressor.html>)

```
In [ ]: rf_random.best_params_
```

```
Out[ ]: {'n_estimators': 300,
    'min_samples_split': 5,
    'min_samples_leaf': 1,
    'max_features': 'sqrt',
    'max_depth': 30}
```

```
In [ ]: best_rf = rf_random.best_estimator_
test_models(X_train, X_test, y_train, y_test, best_rf, "Random Forest (Tunned)", ensemble_bagging_methods)
```

	Model	Test Sc.	Train Sc.	MAE	MSE	RMSE	RMSE CV
0	Random Forest (Tunned)	0.623237	0.819132	0.299825	0.178242	0.422187	0.441452

Boosting Ensemble Methods

Gradient Boosting

```
In [ ]: gbr = GradientBoostingRegressor()
gbr.fit(X_train,y_train)
test_models(X_train, X_test, y_train, y_test, gbr, "Gradient Boosting", ensemble_boosting_methods)
```

	Model	Test Sc.	Train Sc.	MAE	MSE	RMSE	RMSE CV
0	Gradient Boosting	0.592386	0.590188	0.315598	0.192837	0.439132	0.453508

Hyper-parameter Tuning

```
In [ ]: from pprint import pprint
print('Parameters currently in use: \n')
pprint(gbr.get_params())
```

Parameters currently in use:

```
{'alpha': 0.9,
 'ccp_alpha': 0.0,
 'criterion': 'friedman_mse',
 'init': None,
 'learning_rate': 0.1,
 'loss': 'squared_error',
 'max_depth': 3,
 'max_features': None,
 'max_leaf_nodes': None,
 'min_impurity_decrease': 0.0,
 'min_samples_leaf': 1,
 'min_samples_split': 2,
 'min_weight_fraction_leaf': 0.0,
 'n_estimators': 100,
 'n_iter_no_change': None,
 'random_state': None,
 'subsample': 1.0,
 'tol': 0.0001,
 'validation_fraction': 0.1,
 'verbose': 0,
 'warm_start': False}
```

```
In [ ]: param_grid = {
    "learning_rate": [0.01, 0.05, 0.1, 0.2, 0.5],
    "alpha": [0.5, 0.9, 1.2, 1.5],
    "subsample": [0.6, 0.8, 1.0],
    "max_depth": [5, 10, 15, 20, 30],
    "n_estimators": [100, 150, 200, 300],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 5],
    'max_features': ['auto', 'sqrt', 'log2']
}
gbr_random = RandomizedSearchCV(estimator = gbr, param_distributions = param_g
rid,
                                n_iter = 10, cv = 3, verbose=2, random_state=4
2,
                                n_jobs = -1)
```

```
In [ ]: # Best parameters selected.
gbr_random.fit(X_train, y_train)
gbr_random.best_params_
```

Fitting 3 folds for each of 10 candidates, totalling 30 fits

```
Out[ ]: {'subsample': 1.0,
        'n_estimators': 100,
        'min_samples_split': 10,
        'min_samples_leaf': 5,
        'max_features': 'sqrt',
        'max_depth': 10,
        'learning_rate': 0.2,
        'alpha': 0.5}
```

```
In [ ]: best_gbr = gbr_random.best_estimator_
test_models(X_train, X_test, y_train, y_test, best_gbr, "Gradient Boosting (Tu
nned)", ensemble_boosting_methods)
```

	Model	Test Sc.	Train Sc.	MAE	MSE	RMSE	RMSE CV
0	Gradient Boosting (Tunned)	0.619605	0.693612	0.304394	0.17996	0.424217	0.445204

XGBoost

```
In [ ]: xgbr = xgb.XGBRegressor(tree_method='gpu_hist', gpu_id=0)
xgbr.fit(X_train,y_train)
test_models(X_train, X_test, y_train, y_test, xgbr, "XGBoost", ensemble_boosti
ng_methods)
```

	Model	Test Sc.	Train Sc.	MAE	MSE	RMSE	RMSE CV
0	XGBoost	0.613338	0.692128	0.305603	0.182925	0.427697	0.444667

Hyper-parameter Tuning

```
In [ ]: from pprint import pprint
        print('Parameters currently in use: \n')
        pprint(xgbr.get_params())
```

Parameters currently in use:

```
{'base_score': None,
 'booster': None,
 'callbacks': None,
 'colsample_bylevel': None,
 'colsample_bynode': None,
 'colsample_bytree': None,
 'device': None,
 'early_stopping_rounds': None,
 'enable_categorical': False,
 'eval_metric': None,
 'feature_types': None,
 'gamma': None,
 'gpu_id': 0,
 'grow_policy': None,
 'importance_type': None,
 'interaction_constraints': None,
 'learning_rate': None,
 'max_bin': None,
 'max_cat_threshold': None,
 'max_cat_to_onehot': None,
 'max_delta_step': None,
 'max_depth': None,
 'max_leaves': None,
 'min_child_weight': None,
 'missing': nan,
 'monotone_constraints': None,
 'multi_strategy': None,
 'n_estimators': None,
 'n_jobs': None,
 'num_parallel_tree': None,
 'objective': 'reg:squarederror',
 'random_state': None,
 'reg_alpha': None,
 'reg_lambda': None,
 'sampling_method': None,
 'scale_pos_weight': None,
 'subsample': None,
 'tree_method': 'gpu_hist',
 'validate_parameters': None,
 'verbosity': None}
```

```
In [ ]: param_grid = {
    "learning_rate": [0.01, 0.05, 0.1, 0.2, 0.5],
    "colsample_bytree": [0.6, 0.8, 1.0],
    "subsample": [0.6, 0.8, 1.0],
    "max_depth": [10, 15, 20, 30],
    "n_estimators": [100, 150, 200, 300],
    "reg_lambda": [1, 1.5, 2],
    "gamma": [0, 0.1, 0.3],
    'min_child_weight': [1, 3, 5],
}
xgbr_random = RandomizedSearchCV(estimator = xgbr, param_distributions = param
_grid,
                                n_iter = 10, cv = 3, verbose=2, random_state=9
1,
                                n_jobs = -1)
```

```
In [ ]: xgbr_random.fit(X_train, y_train)
xgbr_random.best_params_
```

Fitting 3 folds for each of 10 candidates, totalling 30 fits

```
Out[ ]: {'subsample': 0.6,
'learning_rate': 0.01,
'gamma': 0,
'colsample_bytree': 0.6,
'reg_lambda': 1.5,
'n_estimators': 300,
'min_child_weight': 5,
'max_depth': 15}
```

```
In [ ]: best_xgb = xgbr_random.best_estimator_
test_models(X_train, X_test, y_train, y_test, best_xgb, "XGBoost (Tunned)", en
semble_boosting_methods)
```

	Model	Test Sc.	Train Sc.	MAE	MSE	RMSE	RMSE CV
0	XGBoost (Tunned)	0.622568	0.765177	0.301494	0.178558	0.422562	0.440758

Comparison of Boosting Methods

```
In [ ]: print("Green Highlights the best model, and Red highlights the worst!!!")
compare_models(pd.DataFrame(ensemble_boosting_methods),30)
```

Green Highlights the best model, and Red highlights the worst!!!

	Model	Test Sc.	Train Sc.	MAE	MSE	RMSE	RMSE CV
0	Gradient Boosting	0.592386	0.590188	0.315598	0.192837	0.439132	0.453508
1	Gradient Boosting (Tunned)	0.619605	0.693612	0.304394	0.179960	0.424217	0.445204
2	XGBoost	0.613338	0.692128	0.305603	0.182925	0.427697	0.444667
3	XGBoost (Tunned)	0.622568	0.765177	0.301494	0.178558	0.422562	0.440758

Comparison of Ensemble methods

```
In [ ]: print("Green Highlights the best model, and Red highlights the worst")
compare_models(pd.DataFrame(ensemble_bagging_methods+ensemble_boosting_methods),45)
```

Green Highlights the best model, and Red highlights the worst

	Model	Test Sc.	Train Sc.	MAE	MSE	RMSE	RMSE CV
0	Random Forest (Base)	0.610250	0.942988	0.306271	0.184386	0.429402	0.447824
1	Random Forest (Tunned)	0.623237	0.819132	0.299825	0.178242	0.422187	0.441452
2	Gradient Boosting	0.592386	0.590188	0.315598	0.192837	0.439132	0.453508
3	Gradient Boosting (Tunned)	0.619605	0.693612	0.304394	0.179960	0.424217	0.445204
4	XGBoost	0.613338	0.692128	0.305603	0.182925	0.427697	0.444667
5	XGBoost (Tunned)	0.622568	0.765177	0.301494	0.178558	0.422562	0.440758

Comparing all models together

```
In [ ]: compare_models(pd.DataFrame(classic_ml_models+ensemble_bagging_methods+ensemble_boosting_methods),60)
```

	Model	Test Sc.	Train Sc.	MAE	MSE	RMSE	RMSE CV
0	Ridge Regression	0.550181	0.537237	0.334657	0.212804	0.461306	0.478069
1	Lasso Regression	0.549930	0.537223	0.334679	0.212923	0.461435	0.477871
2	Decision Tree	0.571749	0.600147	0.326018	0.202600	0.450111	0.467157
3	Random Forest (Base)	0.610250	0.942988	0.306271	0.184386	0.429402	0.447824
4	Random Forest (Tunned)	0.623237	0.819132	0.299825	0.178242	0.422187	0.441452
5	Gradient Boosting	0.592386	0.590188	0.315598	0.192837	0.439132	0.453508
6	Gradient Boosting (Tunned)	0.619605	0.693612	0.304394	0.179960	0.424217	0.445204
7	XGBoost	0.613338	0.692128	0.305603	0.182925	0.427697	0.444667
8	XGBoost (Tunned)	0.622568	0.765177	0.301494	0.178558	0.422562	0.440758

Comments

Of all the evaluated models and methods, the hyper-tuned XGBoost model gave the best results and the Lasso regression gave the worst results. The ensemble methods clearly outperform the classical machine learning models for Airbnb price prediction problem because it helps in the reduction in the variance component of prediction errors made by the contributing models. Moreover, the hyper-tuning further improved the results, which can be further enhanced using a grid search instead of a random search.


Neural Networks


```
In [ ]: from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Dense, Dropout, BatchNormalization
        from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint


ann_model = Sequential([
    Dense(128, activation='relu', input_shape=(X_train.shape[1],)),
    BatchNormalization(),
    Dropout(0.3),
    Dense(64, activation='relu'),
    BatchNormalization(),
    Dropout(0.2),
    Dense(32, activation='relu'),
    Dense(1)
])


early_stopping = EarlyStopping(monitor='val_loss', patience=10, restore_best_weights=True)
model_checkpoint = ModelCheckpoint('best_model.keras', monitor='val_loss', save_best_only=True)
```


```
In [ ]: ann_model.compile(optimizer='adam', loss='mean_squared_error', metrics=['mae'])
        history = ann_model.fit(
            X_train, y_train,
            epochs=100,
            batch_size=32,
            validation_split=0.2,
            callbacks=[early_stopping, model_checkpoint],
            verbose=1
        )
```


Epoch 1/100
978/978  9s 5ms/step - loss: 3.7063 - mae: 1.3502 - val_loss: 0.4399 - val_mae: 0.4506


Epoch 2/100
978/978  5s 2ms/step - loss: 0.5346 - mae: 0.5599 - val_loss: 0.3246 - val_mae: 0.3970


Epoch 3/100
978/978  2s 2ms/step - loss: 0.4134 - mae: 0.4875 - val_loss: 0.3160 - val_mae: 0.3900


Epoch 4/100
978/978  3s 2ms/step - loss: 0.3514 - mae: 0.4465 - val_loss: 0.2996 - val_mae: 0.3684


Epoch 5/100
978/978  3s 3ms/step - loss: 0.3204 - mae: 0.4192 - val_loss: 0.2877 - val_mae: 0.3701


Epoch 6/100
978/978  3s 3ms/step - loss: 0.2830 - mae: 0.3938 - val_loss: 0.2690 - val_mae: 0.3554


Epoch 7/100
978/978  5s 3ms/step - loss: 0.2747 - mae: 0.3808 - val_loss: 0.2536 - val_mae: 0.3527


Epoch 8/100
978/978  2s 2ms/step - loss: 0.2502 - mae: 0.3650 - val_loss: 0.2665 - val_mae: 0.3644


Epoch 9/100
978/978  2s 2ms/step - loss: 0.2516 - mae: 0.3665 - val_loss: 0.2550 - val_mae: 0.3540


Epoch 10/100
978/978  2s 2ms/step - loss: 0.2356 - mae: 0.3564 - val_loss: 0.2465 - val_mae: 0.3489


Epoch 11/100
978/978  3s 3ms/step - loss: 0.2398 - mae: 0.3545 - val_loss: 0.2422 - val_mae: 0.3436


Epoch 12/100
978/978  2s 2ms/step - loss: 0.2357 - mae: 0.3524 - val_loss: 0.2541 - val_mae: 0.3573


Epoch 13/100
978/978  2s 2ms/step - loss: 0.2370 - mae: 0.3485 - val_loss: 0.2675 - val_mae: 0.3474


Epoch 14/100
978/978  2s 2ms/step - loss: 0.2290 - mae: 0.3456 - val_loss: 0.2383 - val_mae: 0.3371

Epoch 15/100
978/978  3s 2ms/step - loss: 0.2255 - mae: 0.3429 - val_loss: 0.2491 - val_mae: 0.3377

Epoch 16/100
978/978  3s 3ms/step - loss: 0.2289 - mae: 0.3448 - val_loss: 0.2502 - val_mae: 0.3510

Epoch 17/100
978/978  3s 3ms/step - loss: 0.2255 - mae: 0.3398 - val_loss: 0.2452 - val_mae: 0.3387

Epoch 18/100
978/978  4s 2ms/step - loss: 0.2244 - mae: 0.3410 - val_loss: 0.2424 - val_mae: 0.3392

Epoch 19/100
978/978  2s 2ms/step - loss: 0.2321 - mae: 0.3432 - val_loss: 0.2409 - val_mae: 0.3397

Epoch 20/100

978/978 ————— 3s 2ms/step - loss: 0.2172 - mae: 0.3363 - val_loss: 0.2430 - val_mae: 0.3365

Epoch 21/100

978/978 ————— 3s 3ms/step - loss: 0.2130 - mae: 0.3345 - val_loss: 0.2410 - val_mae: 0.3378

Epoch 22/100

978/978 ————— 5s 2ms/step - loss: 0.2216 - mae: 0.3383 - val_loss: 0.2405 - val_mae: 0.3329

Epoch 23/100

978/978 ————— 2s 2ms/step - loss: 0.2206 - mae: 0.3358 - val_loss: 0.2397 - val_mae: 0.3361

Epoch 24/100

978/978 ————— 2s 2ms/step - loss: 0.2076 - mae: 0.3290 - val_loss: 0.2386 - val_mae: 0.3326

```
In [ ]: # Evaluate
ann_pred = ann_model.predict(X_test)
print("\nArtificial Neural Network Results:")
print(f"MAE: {mean_absolute_error(y_test, ann_pred):.4f}")
print(f"RMSE: {np.sqrt(mean_squared_error(y_test, ann_pred)):.4f}")
print(f"R2 Score: {r2_score(y_test, ann_pred):.4f}")
```

306/306 ————— 1s 3ms/step


Artificial Neural Network Results:

MAE: 0.3298

RMSE: 0.4613

R2 Score: 0.5501


```
In [ ]: plt.figure(figsize=(10,6))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()
```



Hyperparameter Tunning

```
In [ ]: !pip install scikeras
```

```
Collecting scikeras
  Downloading scikeras-0.13.0-py3-none-any.whl.metadata (3.1 kB)
Requirement already satisfied: keras>=3.2.0 in /usr/local/lib/python3.10/dist-packages (from scikeras) (3.5.0)
Requirement already satisfied: scikit-learn>=1.4.2 in /usr/local/lib/python3.10/dist-packages (from scikeras) (1.5.2)
Requirement already satisfied: absl-py in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->scikeras) (1.4.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->scikeras) (1.26.4)
Requirement already satisfied: rich in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->scikeras) (13.9.4)
Requirement already satisfied: namex in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->scikeras) (0.0.8)
Requirement already satisfied: h5py in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->scikeras) (3.12.1)
Requirement already satisfied: optree in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->scikeras) (0.13.1)
Requirement already satisfied: ml-dtypes in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->scikeras) (0.4.1)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->scikeras) (24.2)
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=1.4.2->scikeras) (1.13.1)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=1.4.2->scikeras) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn>=1.4.2->scikeras) (3.5.0)
Requirement already satisfied: typing-extensions>=4.5.0 in /usr/local/lib/python3.10/dist-packages (from optree->keras>=3.2.0->scikeras) (4.12.2)
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.10/dist-packages (from rich->keras>=3.2.0->scikeras) (3.0.0)
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.10/dist-packages (from rich->keras>=3.2.0->scikeras) (2.18.0)
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.10/dist-packages (from markdown-it-py>=2.2.0->rich->keras>=3.2.0->scikeras) (0.1.2)
Downloading scikeras-0.13.0-py3-none-any.whl (26 kB)
Installing collected packages: scikeras
Successfully installed scikeras-0.13.0
```

Hyperparameter Tunning

```

In [ ]: from scipy.stats import randint, uniform
        from tensorflow.keras.optimizers import Adam
        from sklearn.model_selection import RandomizedSearchCV
        from scikeras.wrappers import KerasRegressor

def create_model(neurons1=128, neurons2=64, neurons3=32, dropout1=0.3, dropout
2=0.2, learning_rate=0.001):
    model = Sequential([
        Dense(neurons1, activation='relu', input_shape=(X_train.shape[1],)),
        BatchNormalization(),
        Dropout(dropout1),
        Dense(neurons2, activation='relu'),
        BatchNormalization(),
        Dropout(dropout2),
        Dense(neurons3, activation='relu'),
        Dense(1)
    ])
    model.compile(optimizer=Adam(learning_rate=learning_rate), loss='mean_squa
red_error')
    return model

param_dist = {
    'model__neurons1': randint(64, 256),
    'model__neurons2': randint(32, 128),
    'model__neurons3': randint(16, 64),
    'model__dropout1': uniform(0.1, 0.3),
    'model__dropout2': uniform(0.1, 0.3),
    'model__learning_rate': uniform(0.001, 0.01),
    'batch_size': randint(32, 64),
    'epochs': randint(50, 100)
}

model = KerasRegressor(model=create_model, verbose=0)
random_search = RandomizedSearchCV(
    estimator=model,
    param_distributions=param_dist,
    n_iter=5,
    cv=3,
    verbose=2,
    n_jobs=1
)

random_search.fit(X_train, y_train)

print("Best parameters:", random_search.best_params_)
print("Best score:", random_search.best_score_)

# Use best parameters to create final model
best_model = create_model(**{k.replace('model__', ''): v for k, v in random_se
arch.best_params_.items()
                           if k.startswith('model__')})

```

```
# Train final model
final_history = best_model.fit(
    X_train, y_train,
    epochs=random_search.best_params_['epochs'],
    batch_size=random_search.best_params_['batch_size'],
    verbose=1
)
```

Fitting 3 folds for each of 5 candidates, totalling 15 fits

```
[CV] END batch_size=41, epochs=87, model__dropout1=0.1485219233572097, model__dropout2=0.3370253488136214, model__learning_rate=0.0041809010854051715, model__neurons1=109, model__neurons2=70, model__neurons3=32; total time= 2.4min
[CV] END batch_size=41, epochs=87, model__dropout1=0.1485219233572097, model__dropout2=0.3370253488136214, model__learning_rate=0.0041809010854051715, model__neurons1=109, model__neurons2=70, model__neurons3=32; total time= 1.9min
[CV] END batch_size=41, epochs=87, model__dropout1=0.1485219233572097, model__dropout2=0.3370253488136214, model__learning_rate=0.0041809010854051715, model__neurons1=109, model__neurons2=70, model__neurons3=32; total time= 2.0min
[CV] END batch_size=36, epochs=99, model__dropout1=0.14294668256567616, model__dropout2=0.11323830608935137, model__learning_rate=0.004391039777312497, model__neurons1=237, model__neurons2=113, model__neurons3=17; total time= 2.4min
[CV] END batch_size=36, epochs=99, model__dropout1=0.14294668256567616, model__dropout2=0.11323830608935137, model__learning_rate=0.004391039777312497, model__neurons1=237, model__neurons2=113, model__neurons3=17; total time= 2.5min
[CV] END batch_size=36, epochs=99, model__dropout1=0.14294668256567616, model__dropout2=0.11323830608935137, model__learning_rate=0.004391039777312497, model__neurons1=237, model__neurons2=113, model__neurons3=17; total time= 2.5min
[CV] END batch_size=36, epochs=75, model__dropout1=0.3413943528837057, model__dropout2=0.1815656253644331, model__learning_rate=0.005807098523600376, model__neurons1=73, model__neurons2=73, model__neurons3=33; total time= 1.9min
[CV] END batch_size=36, epochs=75, model__dropout1=0.3413943528837057, model__dropout2=0.1815656253644331, model__learning_rate=0.005807098523600376, model__neurons1=73, model__neurons2=73, model__neurons3=33; total time= 1.9min
[CV] END batch_size=36, epochs=75, model__dropout1=0.3413943528837057, model__dropout2=0.1815656253644331, model__learning_rate=0.005807098523600376, model__neurons1=73, model__neurons2=73, model__neurons3=33; total time= 1.8min
[CV] END batch_size=59, epochs=80, model__dropout1=0.2467979824448353, model__dropout2=0.33753766195123713, model__learning_rate=0.007263643743275193, model__neurons1=84, model__neurons2=53, model__neurons3=59; total time= 1.5min
[CV] END batch_size=59, epochs=80, model__dropout1=0.2467979824448353, model__dropout2=0.33753766195123713, model__learning_rate=0.007263643743275193, model__neurons1=84, model__neurons2=53, model__neurons3=59; total time= 1.4min
[CV] END batch_size=59, epochs=80, model__dropout1=0.2467979824448353, model__dropout2=0.33753766195123713, model__learning_rate=0.007263643743275193, model__neurons1=84, model__neurons2=53, model__neurons3=59; total time= 1.4min
[CV] END batch_size=48, epochs=74, model__dropout1=0.26588772235915503, model__dropout2=0.3011453109290326, model__learning_rate=0.0018375896254714634, model__neurons1=78, model__neurons2=51, model__neurons3=24; total time= 2.0min
[CV] END batch_size=48, epochs=74, model__dropout1=0.26588772235915503, model__dropout2=0.3011453109290326, model__learning_rate=0.0018375896254714634, model__neurons1=78, model__neurons2=51, model__neurons3=24; total time= 1.6min
[CV] END batch_size=48, epochs=74, model__dropout1=0.26588772235915503, model__dropout2=0.3011453109290326, model__learning_rate=0.0018375896254714634, model__neurons1=78, model__neurons2=51, model__neurons3=24; total time= 1.7min
Best parameters: {'batch_size': 36, 'epochs': 75, 'model__dropout1': 0.3413943528837057, 'model__dropout2': 0.1815656253644331, 'model__learning_rate': 0.005807098523600376, 'model__neurons1': 73, 'model__neurons2': 73, 'model__neurons3': 33}
```





























Best score: 0.5546596551886979

Epoch 1/75

1087/1087  7s 4ms/step - loss: 1.5131

Epoch 2/75

1087/1087	3s	2ms/step	-	loss: 0.3268
Epoch 3/75				
1087/1087	2s	2ms/step	-	loss: 0.2704
Epoch 4/75				
1087/1087	3s	2ms/step	-	loss: 0.2546
Epoch 5/75				
1087/1087	2s	2ms/step	-	loss: 0.2579
Epoch 6/75				
1087/1087	2s	2ms/step	-	loss: 0.2544
Epoch 7/75				
1087/1087	2s	2ms/step	-	loss: 0.2476
Epoch 8/75				
1087/1087	3s	3ms/step	-	loss: 0.2399
Epoch 9/75				
1087/1087	2s	2ms/step	-	loss: 0.2459
Epoch 10/75				
1087/1087	2s	2ms/step	-	loss: 0.2347
Epoch 11/75				
1087/1087	2s	2ms/step	-	loss: 0.2314
Epoch 12/75				
1087/1087	3s	2ms/step	-	loss: 0.2352
Epoch 13/75				
1087/1087	3s	2ms/step	-	loss: 0.2350
Epoch 14/75				
1087/1087	3s	3ms/step	-	loss: 0.2275
Epoch 15/75				
1087/1087	4s	2ms/step	-	loss: 0.2249
Epoch 16/75				
1087/1087	2s	2ms/step	-	loss: 0.2366
Epoch 17/75				
1087/1087	3s	2ms/step	-	loss: 0.2255
Epoch 18/75				
1087/1087	3s	2ms/step	-	loss: 0.2255
Epoch 19/75				
1087/1087	3s	2ms/step	-	loss: 0.2272
Epoch 20/75				
1087/1087	2s	2ms/step	-	loss: 0.2209
Epoch 21/75				
1087/1087	3s	2ms/step	-	loss: 0.2222
Epoch 22/75				
1087/1087	3s	2ms/step	-	loss: 0.2219
Epoch 23/75				
1087/1087	2s	2ms/step	-	loss: 0.2313
Epoch 24/75				
1087/1087	3s	2ms/step	-	loss: 0.2240
Epoch 25/75				
1087/1087	3s	2ms/step	-	loss: 0.2187
Epoch 26/75				
1087/1087	2s	2ms/step	-	loss: 0.2229
Epoch 27/75				
1087/1087	2s	2ms/step	-	loss: 0.2219
Epoch 28/75				
1087/1087	2s	2ms/step	-	loss: 0.2187
Epoch 29/75				
1087/1087	2s	2ms/step	-	loss: 0.2157
Epoch 30/75				
1087/1087	2s	2ms/step	-	loss: 0.2207

Epoch 31/75	
1087/1087 	3s 3ms/step - loss: 0.2252
Epoch 32/75	
1087/1087 	4s 2ms/step - loss: 0.2258
Epoch 33/75	
1087/1087 	2s 2ms/step - loss: 0.2182
Epoch 34/75	
1087/1087 	3s 2ms/step - loss: 0.2164
Epoch 35/75	
1087/1087 	3s 2ms/step - loss: 0.2176
Epoch 36/75	
1087/1087 	5s 2ms/step - loss: 0.2197
Epoch 37/75	
1087/1087 	2s 2ms/step - loss: 0.2194
Epoch 38/75	
1087/1087 	2s 2ms/step - loss: 0.2081
Epoch 39/75	
1087/1087 	3s 2ms/step - loss: 0.2136
Epoch 40/75	
1087/1087 	2s 2ms/step - loss: 0.2175
Epoch 41/75	
1087/1087 	3s 2ms/step - loss: 0.2158
Epoch 42/75	
1087/1087 	3s 3ms/step - loss: 0.2108
Epoch 43/75	
1087/1087 	4s 2ms/step - loss: 0.2157
Epoch 44/75	
1087/1087 	3s 2ms/step - loss: 0.2113
Epoch 45/75	
1087/1087 	2s 2ms/step - loss: 0.2172
Epoch 46/75	
1087/1087 	3s 2ms/step - loss: 0.2156
Epoch 47/75	
1087/1087 	4s 2ms/step - loss: 0.2185
Epoch 48/75	
1087/1087 	3s 2ms/step - loss: 0.2093
Epoch 49/75	
1087/1087 	2s 2ms/step - loss: 0.2142
Epoch 50/75	
1087/1087 	2s 2ms/step - loss: 0.2201
Epoch 51/75	
1087/1087 	3s 2ms/step - loss: 0.2160
Epoch 52/75	
1087/1087 	4s 2ms/step - loss: 0.2205
Epoch 53/75	
1087/1087 	3s 2ms/step - loss: 0.2199
Epoch 54/75	
1087/1087 	3s 2ms/step - loss: 0.2114
Epoch 55/75	
1087/1087 	2s 2ms/step - loss: 0.2114
Epoch 56/75	
1087/1087 	3s 3ms/step - loss: 0.2090
Epoch 57/75	
1087/1087 	2s 2ms/step - loss: 0.2086
Epoch 58/75	
1087/1087 	2s 2ms/step - loss: 0.2147
Epoch 59/75	

```

1087/1087 ————— 3s 2ms/step - loss: 0.2133
Epoch 60/75
1087/1087 ————— 3s 2ms/step - loss: 0.2244
Epoch 61/75
1087/1087 ————— 2s 2ms/step - loss: 0.2143
Epoch 62/75
1087/1087 ————— 3s 2ms/step - loss: 0.2103
Epoch 63/75
1087/1087 ————— 4s 2ms/step - loss: 0.2067
Epoch 64/75
1087/1087 ————— 2s 2ms/step - loss: 0.2178
Epoch 65/75
1087/1087 ————— 2s 2ms/step - loss: 0.2087
Epoch 66/75
1087/1087 ————— 3s 2ms/step - loss: 0.2157
Epoch 67/75
1087/1087 ————— 3s 3ms/step - loss: 0.2113
Epoch 68/75
1087/1087 ————— 2s 2ms/step - loss: 0.2145
Epoch 69/75
1087/1087 ————— 2s 2ms/step - loss: 0.2123
Epoch 70/75
1087/1087 ————— 3s 2ms/step - loss: 0.2152
Epoch 71/75
1087/1087 ————— 3s 2ms/step - loss: 0.2142
Epoch 72/75
1087/1087 ————— 3s 2ms/step - loss: 0.2125
Epoch 73/75
1087/1087 ————— 3s 2ms/step - loss: 0.2145
Epoch 74/75
1087/1087 ————— 4s 2ms/step - loss: 0.2136
Epoch 75/75
1087/1087 ————— 3s 2ms/step - loss: 0.2150

```

```

In [ ]: y_pred = best_model.predict(X_test)

print('Hypertuned Model Test Results:')
print(f'MAE: {mean_absolute_error(y_test, y_pred):.4f}')
print(f'RMSE: {np.sqrt(mean_squared_error(y_test, y_pred)):.4f}')
print(f'R2 Score: {r2_score(y_test, y_pred):.4f}')

```

```

306/306 ————— 2s 5ms/step
Hypertuned Model Test Results:
MAE: 0.3203
RMSE: 0.4480
R2 Score: 0.5757

```

While the prediction enhanced after hyperparameter tuning, our artificial neural network is not efficient as the Random Forest model.